

SUMMER INTERNSHIP REPORT

PNEUMONIA DETECTION SYSTEM

SUBMITTED BY Yashas Raj R Enrollment No.: 24162121033	INDUSTRY PARTNER Akshar AI Ltd Summer Internship 2026
---	---

Guided by
Prof. Dhaval Sathawara

In partial fulfilment for the award of the degree of
BACHELOR OF TECHNOLOGY
In
Computer Science Engineering
Faculty of Engineering and Technology

Certificate

This is to certify that the Summer Internship Report entitled
"Pneumonia Detection System Using Deep Learning"
submitted by

Yashas Raj R (Enrollment No. 24162121033)

in partial fulfilment of the requirements for the award of the degree of
Bachelor of Technology in Computer Science and Engineering (Big Data Analytics) at **ICT, Ganpat University, Mehsana, Gujarat**, is a bona fide record of the work carried out during the Summer Internship at **Akshar AI Ltd.** under the guidance and supervision of the industry mentor and faculty guide.

The work presented in this report is original and has been completed by the student during the internship period. It demonstrates the application of Artificial Intelligence and Deep Learning techniques for the detection of pneumonia from chest X-ray images through the development of a web-based diagnostic support system. The project involved data collection, preprocessing, model development, testing, deployment, and integration of the trained model into a user-friendly web application. To the best of our knowledge, the work embodied in this report has not been submitted, either in part or in full, for the award of any other degree, diploma, or similar qualification at this or any other institution.

Industry Mentor
Akshar AI Ltd

Faculty Guide
ICT Ganpat University

Declaration

I hereby declare that the internship report entitled "Pneumonia Detection System " is an authentic, original record of my own work carried out at Akshar AI Ltd under the expert guidance of my industry mentors and the academic direction of my faculty advisors.

This report has been submitted in partial fulfillment of the rigorous requirements for the award of the B.Tech degree in Computer Science and Engineering (BDA) at ICT Ganpat University, Mehsana, Gujarat.

I further declare that the data, methodologies, code snippets, and research work presented in this document have not been submitted elsewhere for the award of any other degree, diploma, fellowship, or other similar titles or prizes. All sources of information and third-party libraries have been duly acknowledged in the References section.

Yashas Raj R

Date: 23 July 2026

Acknowledgement

I would like to express my deepest and most sincere gratitude to all those who provided me with the opportunity, resources, and mentorship necessary to complete this extensive summer internship and the subsequent final project.

First and foremost, I am profoundly grateful to the leadership and engineering teams at Akshar AI Ltd. Providing an intern with the autonomy to tackle complex, real-world problems in medical imaging and computer vision is a rare privilege. The exposure to professional software engineering practices, agile methodologies, and production-grade deployments has been invaluable to my career trajectory.

I wish to extend special thanks to my assigned industry mentors. Their continuous support, technical guidance, code reviews, and extraordinary patience throughout the 8-week program were the primary catalysts for my success. Their deep insights into deep learning optimization, the nuances of transfer learning, and the intricacies of building robust Flask APIs elevated my work from academic prototypes to deployable software.

I also extend my sincere appreciation to the faculty, administration, and placement cell at ICT Ganpat University. The foundational theoretical knowledge imparted during my coursework provided the mathematical and computational bedrock required to understand and execute these advanced artificial intelligence algorithms.

Finally, I want to thank my family and friends for their unwavering support, encouragement, and understanding during the long hours of research, debugging, and model training that characterized this intensive learning period.

Abstract

Pneumonia is a serious respiratory disease that affects millions of people worldwide and remains one of the leading causes of illness and mortality, particularly among children and older adults. Early and accurate diagnosis is essential for timely treatment; however, manual interpretation of chest X-ray images can be time-consuming and depends heavily on the expertise of radiologists. Recent advancements in Artificial Intelligence (AI) and Deep Learning have demonstrated significant potential in assisting medical professionals by automating the analysis of medical images and improving diagnostic efficiency.

This internship project presents the design and development of a Pneumonia Detection System Using Deep Learning, which classifies chest X-ray images as Normal or Pneumonia. The project involved the complete software development lifecycle, including requirement analysis, dataset collection, data preprocessing, exploratory data analysis, model development, testing, deployment, and documentation. A publicly available Chest X-ray Images (Pneumonia) dataset from Kaggle was used to train and evaluate the models. Image preprocessing techniques such as resizing, normalization, and data validation were applied to prepare the dataset for deep learning.

Three deep learning models—Custom Convolutional Neural Network (CNN), MobileNetV2, and ResNet50—were implemented and compared based on their classification performance. Among these, MobileNetV2 demonstrated the best balance between prediction accuracy, computational efficiency, and inference speed, making it the most suitable model for deployment. The selected model was integrated into a web application developed using React.js for the frontend and Flask for the backend. The application allows authenticated users to upload chest X-ray images and receive real-time predictions through an intuitive and responsive user interface.

The developed system was tested using standard performance evaluation metrics, including accuracy, precision, recall, F1-score, and confusion matrix, to validate its effectiveness. Functional and integration testing confirmed the reliability of the application and the seamless communication between the frontend, backend, and deep learning model. The project demonstrates how AI-assisted diagnostic systems can support healthcare professionals by providing fast and reliable preliminary screening of pneumonia cases.

Overall, the internship provided practical exposure to deep learning, medical image analysis, full-stack web development, model deployment, and software engineering practices. The developed Pneumonia Detection System serves as a scalable foundation for future enhancements, including multi-class disease detection, Explainable AI (XAI), cloud deployment, and integration with hospital information systems.

Keywords: *Artificial Intelligence, Deep Learning, Pneumonia Detection, Chest X-ray, Convolutional Neural Network, MobileNetV2, ResNet50, React, Flask, Medical Image Classification.*

Table of Contents

CHAPTER 1: INTRODUCTION

- 1.1 Project Overview
- 1.2 Background of the Study
- 1.3 Problem Statement
- 1.4 Project Objectives
- 1.5 Scope of the Project
- 1.6 Project Deliverables
- 1.7 Report Organization

CHAPTER 2: REQUIREMENT ANALYSIS

- 2.1 Domain Study
- 2.2 Existing System Analysis
- 2.3 Proposed System
- 2.4 Functional Requirements
- 2.5 Non-Functional Requirements
- 2.6 Tools and Technologies Used

CHAPTER 3: DATA COLLECTION AND UNDERSTANDING

- 3.1 Data Source
- 3.2 Data Description
- 3.3 Data Structure
- 3.4 Data Dictionary
- 3.5 Data Quality Assessment

CHAPTER 4: DATA PREPROCESSING

- 4.1 Data Cleaning
- 4.2 Handling Missing Values
- 4.3 Duplicate Removal
- 4.4 Data Transformation
- 4.5 Data Validation
- 4.6 Final Prepared Dataset

CHAPTER 5: DATA ANALYSIS AND MODELING

- 5.1 Exploratory Data Analysis (EDA)
- 5.2 Data Relationships
- 5.3 Key Metrics Identification
- 5.4 Business Rules and Logic
- 5.5 Data Modeling Approach

CHAPTER 6: DASHBOARD DESIGN AND DEVELOPMENT

- 6.1 Dashboard Planning
- 6.2 Data Model Design
- 6.3 Dashboard Components

6.4	Visualizations Used
6.5	Interactive Features
6.6	User Interface Design
CHAPTER 7: RESULTS AND INSIGHTS	
7.1	Dashboard Output
7.2	Key Findings
7.3	Trend Analysis
7.4	Comparative Analysis
7.5	Generated Insights
CHAPTER 8: TESTING AND VALIDATION	
8.1	Testing Strategy
8.2	Test Cases
8.3	Validation Process
8.4	Performance Evaluation
8.5	Limitations
CHAPTER 9: DEPLOYMENT	
9.1	Deployment Process
9.2	Publishing and Sharing
9.3	User Access and Usage
9.4	Documentation
CHAPTER 10: CONCLUSION AND FUTURE SCOPE	
10.1	Project Summary
10.2	Objectives Achieved
10.3	Learning Outcomes
10.4	Future Enhancements
REFERENCES	

CHAPTER1

INTRODUCTION

1.1 Project Overview

The rapid advancement of Artificial Intelligence (AI) and Deep Learning (DL) technologies has transformed numerous industries, with healthcare emerging as one of the most impactful domains. Modern healthcare systems generate vast amounts of medical data every day in the form of laboratory reports, electronic health records, medical images, and diagnostic reports. Extracting meaningful information from these data sources is essential for improving disease diagnosis, treatment planning, and patient care. Among the various applications of AI in healthcare, medical image analysis has gained considerable attention due to its ability to assist clinicians in detecting diseases quickly and accurately.

Medical imaging techniques such as X-ray, Computed Tomography (CT), Magnetic Resonance Imaging (MRI), and Ultrasound play an essential role in diagnosing numerous diseases. Chest X-ray imaging, in particular, is one of the most widely used diagnostic techniques for identifying respiratory disorders because it is inexpensive, non-invasive, and readily available in hospitals and diagnostic centers. Physicians and radiologists analyze these images to identify abnormalities associated with diseases such as pneumonia, tuberculosis, lung cancer, and other pulmonary infections. However, the increasing number of patients requiring radiological examinations places a significant burden on healthcare professionals, making rapid and accurate diagnosis challenging.

Pneumonia is a severe respiratory infection that affects the lungs by causing inflammation in the alveoli (air sacs), which may become filled with fluid or pus. The disease can result from bacterial, viral, or fungal infections and remains one of the leading causes of mortality worldwide, particularly among children under five years of age, elderly individuals, and patients with weakened immune systems. According to global health organizations, early diagnosis and timely treatment are crucial in reducing complications and improving survival rates. Nevertheless, accurate interpretation of chest X-ray images requires skilled radiologists, and healthcare facilities in remote or resource-constrained regions often face shortages of experienced specialists.

Artificial Intelligence offers a promising solution to these challenges. Deep learning models, particularly Convolutional Neural Networks (CNNs), have demonstrated exceptional performance in image classification tasks by automatically learning hierarchical image features without manual feature engineering. Instead of relying solely on handcrafted image descriptors, CNN-based models can learn complex visual patterns directly from training images, making them highly effective for medical image analysis. Transfer learning architectures such as MobileNetV2 and ResNet50 further improve classification performance while significantly reducing training time and computational requirements.

The primary objective of this internship project is to design and develop a Pneumonia Detection System capable of automatically predicting whether a chest X-ray image belongs to a patient with pneumonia or a healthy individual. The project combines artificial intelligence techniques with modern web technologies to create an interactive application that enables users to upload chest X-ray images and receive prediction results within a few seconds. The developed system serves as a clinical decision-support tool rather than replacing medical professionals, assisting healthcare practitioners in preliminary diagnosis and reducing diagnostic workload.

The development of this project followed a structured software engineering approach consisting of several phases. Initially, a publicly available chest X-ray dataset was collected and analyzed to understand its structure and quality. Exploratory Data Analysis (EDA) was then performed to examine class distribution, image characteristics, and dataset consistency. During the preprocessing phase, images were resized, normalized, and prepared for deep learning model training. Multiple neural network architectures, including a custom CNN, MobileNetV2, and ResNet50, were implemented and evaluated to identify the most suitable model for deployment. In addition to model development, the project also focused on building a complete web application capable of providing a seamless user experience. A React-based frontend was developed to enable user authentication, image upload, and result visualization. A Flask backend was implemented to handle server-side operations, communicate with the trained deep learning model, process uploaded images, and return prediction results. The integration of these technologies resulted in a full-stack AI-powered healthcare application that demonstrates the practical application of machine learning concepts in solving real-world medical problems.

Another important aspect of this project is its emphasis on usability and accessibility. Instead of limiting the work to model training, the internship involved integrating the trained model into a functional software system that could potentially be used by healthcare professionals, researchers, and students. The developed application simplifies the prediction process by providing an intuitive graphical interface while maintaining the computational capabilities of the underlying deep learning model.

From an educational perspective, this internship provided extensive exposure to various interdisciplinary domains, including machine learning, deep learning, computer vision, medical image processing, frontend development, backend development, software engineering, model deployment, and version control. Working on an end-to-end AI project significantly enhanced practical understanding of modern software development practices while reinforcing theoretical concepts learned during academic coursework.

Overall, the Pneumonia Detection System represents a comprehensive application of artificial intelligence in healthcare. By integrating deep learning algorithms with an interactive web platform, the project demonstrates how intelligent software systems can assist in improving diagnostic efficiency, reducing workload for healthcare professionals, and supporting early detection of life-threatening diseases. The knowledge and experience gained during the internship provide a strong foundation for future work in artificial intelligence, medical imaging, and healthcare technology.

1.2 Background of the Study

Healthcare has undergone significant technological transformation over the past few decades due to rapid advancements in computing, digital communication, and artificial intelligence. Modern hospitals and healthcare institutions increasingly rely on intelligent systems to improve diagnostic accuracy, automate repetitive tasks, optimize resource utilization, and enhance patient outcomes. Among the various technological innovations, artificial intelligence has emerged as one of the most influential fields because of its ability to analyze large volumes of data, identify complex patterns, and provide predictive insights that support medical decision-making.

Artificial Intelligence refers to the capability of computer systems to perform tasks that normally require human intelligence, including learning, reasoning, problem-solving, decision-making, and pattern recognition. Machine Learning, a subset of AI, enables systems to improve their performance by learning from data instead of being explicitly programmed. Deep Learning, which is a specialized branch of Machine Learning, utilizes multi-layered artificial neural networks to solve complex problems involving image recognition, natural language processing, speech recognition, and autonomous systems.

One of the most successful applications of deep learning is image classification using Convolutional Neural Networks (CNNs). CNNs automatically extract relevant visual features such as edges, textures, shapes, and high-level image representations through multiple convolutional layers. Unlike traditional computer vision techniques that depend on handcrafted features, CNNs learn feature representations directly from data, resulting in significantly improved accuracy across various image recognition tasks.

The healthcare industry has particularly benefited from CNN-based image analysis. Medical images generated through X-rays, CT scans, MRI scans, mammography, and retinal imaging contain valuable diagnostic information that can be analyzed using deep learning algorithms. Researchers worldwide have developed AI systems capable of detecting diseases such as diabetic retinopathy, breast cancer, brain tumors, skin cancer, Alzheimer's disease, and various pulmonary disorders with promising accuracy.

Among respiratory diseases, pneumonia remains one of the most significant global health concerns. Pneumonia is an infection that inflames the air sacs of one or both lungs. These air sacs may become filled with fluid or pus, causing symptoms such as fever, cough, chest pain, breathing difficulty, fatigue, and reduced oxygen levels. Depending on the severity of the infection, pneumonia can become life-threatening, particularly for vulnerable populations including infants, elderly individuals, and immunocompromised patients.

The diagnosis of pneumonia generally involves clinical examination, laboratory testing, and radiological imaging. Chest X-ray examination remains one of the most common diagnostic tools because it provides a clear visualization of lung abnormalities. Radiologists examine these X-ray images to identify regions of opacity, consolidation, or inflammation that indicate pneumonia. However, manual interpretation requires years of medical expertise, and diagnostic accuracy may vary depending on the radiologist's experience, workload, and image quality.

The increasing demand for radiological services has created challenges for healthcare systems worldwide. Many hospitals experience shortages of trained radiologists, resulting in delayed diagnosis and increased workload. In rural and underdeveloped regions, access to experienced specialists may be limited, further increasing the risk of delayed treatment. These challenges have motivated researchers to investigate AI-based diagnostic systems capable of supporting medical professionals by providing rapid preliminary assessments.

Recent developments in transfer learning have accelerated research in medical image analysis. Instead of training large neural networks from scratch, transfer learning utilizes models that have already learned general image features from massive datasets such as ImageNet. These pretrained models can then be fine-tuned using comparatively smaller medical datasets, significantly reducing computational requirements while maintaining high classification

accuracy. Popular transfer learning models such as MobileNetV2, ResNet50, DenseNet121, EfficientNet, and InceptionV3 have demonstrated excellent performance across numerous healthcare applications.

This internship project explores the practical implementation of deep learning techniques for pneumonia detection by comparing multiple neural network architectures and integrating the best-performing model into a user-friendly web application. Beyond developing an accurate classification model, the project emphasizes complete system development, including dataset preprocessing, frontend implementation, backend integration, user authentication, deployment, and documentation.

The significance of this project extends beyond academic learning. The developed Pneumonia Detection System demonstrates how artificial intelligence can assist healthcare professionals by reducing diagnosis time, improving consistency in image interpretation, and providing accessible diagnostic support in environments where expert radiologists may not always be available. Although the system is not intended to replace medical professionals, it illustrates the growing role of intelligent decision-support systems in modern healthcare.

Furthermore, this project aligns with the broader vision of digital healthcare and smart medical technologies, where AI-powered systems complement human expertise to improve healthcare delivery. As artificial intelligence continues to evolve, similar systems are expected to become increasingly integrated into routine clinical practice, supporting physicians in delivering faster, more accurate, and more efficient patient care.

1.3 Problem Statement

Pneumonia is one of the most common and life-threatening respiratory diseases, affecting millions of people every year across the world. It is characterized by inflammation of the lung tissues, primarily the alveoli, due to bacterial, viral, or fungal infections. According to global healthcare reports, pneumonia is responsible for a significant number of hospital admissions and remains one of the leading causes of mortality among children below the age of five, elderly individuals, and patients with weakened immune systems. Early diagnosis plays a critical role in improving treatment outcomes, minimizing complications, and reducing mortality rates.

The conventional diagnosis of pneumonia primarily depends on clinical examination followed by radiological analysis of chest X-ray images. Radiologists examine these X-ray images to identify abnormalities such as infiltrates, consolidations, or opacities that indicate the presence of infection within the lungs. Although chest X-ray examination is considered an effective diagnostic technique, interpreting medical images requires specialized knowledge and considerable clinical experience. Furthermore, image interpretation is often time-consuming, particularly in hospitals handling a large number of patients daily.

Many healthcare institutions, especially those located in rural or economically underdeveloped regions, face a shortage of qualified radiologists. The limited availability of experienced medical professionals often results in delayed diagnosis and treatment. Even in well-equipped hospitals, increasing patient loads may lead to longer waiting times, increased workload for healthcare professionals, and a higher possibility of diagnostic errors due to fatigue or human oversight.

Such delays can significantly affect patient outcomes because pneumonia requires immediate medical intervention to prevent severe complications.

Another challenge associated with manual diagnosis is the subjectivity involved in interpreting chest X-ray images. Different radiologists may occasionally provide different opinions when examining the same image, especially during the early stages of infection where abnormalities are subtle. Variations in image quality, exposure levels, and patient positioning further increase the complexity of diagnosis. Therefore, maintaining consistency and accuracy in manual diagnosis remains a challenging task.

Recent advancements in Artificial Intelligence and Deep Learning have opened new opportunities for addressing these limitations. Convolutional Neural Networks (CNNs) have demonstrated remarkable success in automatically extracting meaningful visual features from images and performing highly accurate image classification. Transfer learning models such as MobileNetV2 and ResNet50 have further improved classification performance while reducing computational complexity and training time. These models have shown promising results in various medical image analysis tasks, including pneumonia detection from chest X-ray images.

Despite these technological advancements, many existing AI models remain limited to research environments and are not integrated into practical software applications that can be easily used by healthcare professionals or researchers. Several published models focus only on improving classification accuracy without addressing system usability, user interaction, deployment, or accessibility. As a result, there remains a gap between AI research and its practical implementation in real-world healthcare systems.

The primary problem addressed in this internship project is the development of an intelligent and user-friendly pneumonia detection system that combines accurate deep learning models with a complete web application. The proposed system aims to simplify the diagnosis process by enabling users to upload chest X-ray images through a web interface and receive automated predictions within a few seconds. By integrating artificial intelligence with modern web technologies, the project seeks to provide a practical decision-support tool capable of assisting healthcare professionals during preliminary diagnosis.

The system is not intended to replace radiologists or medical practitioners. Instead, it serves as an intelligent assistant that helps reduce diagnostic workload, improves consistency, and supports faster decision-making. The successful implementation of such systems has the potential to improve healthcare accessibility, particularly in regions where specialized medical expertise is limited.

1.4 Project Objectives

The primary objective of this internship project is to design, develop, and deploy an intelligent web-based Pneumonia Detection System capable of automatically classifying chest X-ray images as either Normal or Pneumonia using deep learning techniques. The project combines artificial intelligence, computer vision, and full-stack web development to create a practical healthcare application that assists in preliminary diagnosis while providing an intuitive user experience.

To achieve this primary objective, several specific objectives were identified and successfully implemented during the internship period.

The first objective was to study the healthcare domain related to pneumonia diagnosis and understand the importance of chest X-ray imaging in identifying respiratory infections. This involved reviewing the disease, existing diagnostic procedures, challenges associated with manual diagnosis, and recent developments in artificial intelligence for medical image analysis. Understanding the application domain provided the necessary foundation for selecting appropriate technologies and designing the overall system.

The second objective was to collect and analyze a reliable dataset containing chest X-ray images representing both normal and pneumonia cases. Dataset analysis involved examining image quality, dataset structure, class distribution, image dimensions, and overall data consistency. A thorough understanding of the dataset was essential before proceeding to model development.

Another important objective was to preprocess the collected dataset so that it could be effectively used for deep learning model training. Image preprocessing included resizing images to a fixed resolution, normalization of pixel values, organizing images into appropriate directory structures, and preparing training, validation, and testing datasets. Proper preprocessing improves model performance while reducing computational complexity during training.

The project also aimed to perform Exploratory Data Analysis (EDA) to understand the characteristics of the dataset before model development. EDA included visualization of class distribution, examination of sample chest X-ray images, identification of image variations, and assessment of dataset quality. These analyses provided valuable insights that guided subsequent modeling decisions.

A major objective of the internship was the implementation and comparison of multiple deep learning models for pneumonia detection. Instead of relying on a single architecture, different approaches were explored, including a custom Convolutional Neural Network (CNN), MobileNetV2, and ResNet50. The objective was to evaluate each model based on accuracy, computational efficiency, training time, validation performance, and generalization capability to identify the most suitable model for deployment.

Beyond model development, another objective was to integrate the trained deep learning model into a complete web-based software application. This involved designing a React frontend capable of accepting user inputs, uploading medical images, and displaying prediction results in a simple and interactive manner. Simultaneously, a Flask backend was developed to process uploaded images, perform model inference, and return classification results to the frontend.

The internship also focused on implementing user authentication and profile management to enhance application usability and security. Registered users could create accounts, log in securely, upload chest X-ray images, and access prediction services through an organized interface.

Another objective was to evaluate the performance of the developed system using standard machine learning evaluation metrics such as Accuracy, Precision, Recall, F1-Score, and Confusion Matrix. Model evaluation helped identify strengths and limitations while ensuring that the selected model achieved satisfactory classification performance.

Finally, the project aimed to deploy the complete application and document every phase of development, including dataset preparation, model implementation, frontend development, backend integration, testing, deployment, and project management. Proper documentation

ensures that the developed system can be maintained, extended, and understood by future developers and researchers.

The successful completion of these objectives demonstrates the practical application of artificial intelligence in healthcare while strengthening knowledge in deep learning, computer vision, full-stack web development, software engineering, and project deployment.

1.5 Scope of the Project

The scope of the Pneumonia Detection System encompasses the design, development, testing, and deployment of an intelligent software application capable of assisting in the diagnosis of pneumonia using chest X-ray images. The project focuses on applying deep learning techniques to automate image classification while providing a user-friendly web interface that enables seamless interaction between users and the trained AI model.

The primary scope of the project includes binary classification of chest X-ray images into two categories: Normal and Pneumonia. The system is designed to accept chest X-ray images uploaded by authenticated users, preprocess the images according to model requirements, perform prediction using a trained deep learning model, and display the classification result through an interactive web interface.

Another important aspect of the project scope involves the complete machine learning pipeline, beginning with dataset collection and extending through preprocessing, exploratory data analysis, model training, model evaluation, and deployment. Rather than focusing solely on model development, the internship emphasized building an end-to-end software solution capable of real-world implementation.

The scope further includes the development of a React-based frontend for image upload, user authentication, profile management, and prediction visualization. The Flask backend provides API services, handles communication with the trained model, processes uploaded images, and returns prediction results efficiently.

The project also covers comparative evaluation of multiple deep learning architectures to determine the most suitable model for deployment. Performance comparison enables informed model selection while balancing accuracy, computational efficiency, and inference speed.

However, certain limitations define the boundaries of the project. The developed system performs only binary classification and does not diagnose multiple lung diseases simultaneously. It relies solely on chest X-ray images and does not incorporate additional patient information such as age, symptoms, laboratory reports, or medical history. Furthermore, the application is intended as a clinical decision-support tool and should not replace professional medical diagnosis or expert radiological interpretation.

Future enhancements may extend the system to support multiclass disease classification, integration with hospital information systems, explainable AI techniques, cloud deployment, and mobile application support.

1.6 Project Deliverables

The successful completion of an artificial intelligence project is not limited to training a machine learning model; rather, it involves delivering a complete software solution that integrates data processing, model development, user interaction, testing, deployment, and documentation. Throughout this internship, several tangible deliverables were produced, demonstrating both technical implementation and software engineering practices. These deliverables collectively represent the outcome of the project and serve as evidence that the objectives defined at the beginning of the internship were successfully achieved.

One of the primary deliverables of this internship was the creation of a well-organized and preprocessed chest X-ray image dataset suitable for deep learning applications. The original dataset required careful examination before being used for model training. Images were organized according to their respective classes, validated for consistency, resized to a fixed input dimension, and prepared for efficient loading during the training process. This preprocessing stage ensured that the dataset met the requirements of convolutional neural network architectures while maintaining image quality.

Another major deliverable was the Exploratory Data Analysis (EDA) report. Before model development, extensive analysis was performed to understand the characteristics of the dataset. This included studying class distribution, inspecting sample chest X-ray images, identifying variations in image quality, and examining the overall structure of the dataset. The EDA phase provided valuable insights into potential challenges such as class imbalance and image variability, enabling informed decisions during preprocessing and model training.

The most significant deliverable was the development of multiple deep learning models for pneumonia detection. Instead of relying on a single algorithm, different neural network architectures were implemented and evaluated to determine the most effective solution. The project involved:

-
- Development of a Custom Convolutional Neural Network (CNN).
 - Implementation of the MobileNetV2 transfer learning architecture.
 - Implementation of the ResNet50 transfer learning architecture.
 - Comparative evaluation of model performance.
-

Each model was trained using the prepared dataset, validated using separate validation data, and tested to measure classification performance. Performance evaluation included analysis of training accuracy, validation accuracy, training loss, validation loss, confusion matrix, precision, recall, F1-score, and overall classification accuracy. This comparative study enabled the selection of the most suitable model for deployment based on both predictive performance and computational efficiency.

A major software deliverable of the internship was the development of a web-based Pneumonia Detection System. Rather than limiting the project to model training, the internship focused on creating a complete application that users could interact with. The frontend of the application was developed using React, providing a modern, responsive, and user-friendly interface. The frontend allows users to create accounts, log in securely, upload chest X-ray images, and receive prediction results in an organized manner.

The backend was developed using the Flask framework, which serves as the communication layer between the user interface and the trained deep learning model. The backend performs several important functions, including receiving uploaded images, preprocessing them according to model requirements, loading the trained model, generating predictions, and sending prediction results back to the frontend. This integration demonstrates the practical deployment of artificial intelligence models within a real software application.

Another important deliverable was the implementation of user authentication and profile management. The application provides secure registration and login functionality, enabling users to maintain individual accounts. Profile management features allow users to update personal information and securely access prediction services. These functionalities enhance the usability of the system while demonstrating knowledge of full-stack web development principles.

The project also produced a trained deep learning model saved in a deployable format. The trained model can be loaded independently by the Flask backend without requiring retraining each time the application is executed. This significantly improves system performance and enables real-time prediction for uploaded chest X-ray images.

Documentation represents another essential deliverable of the internship. Throughout the project, proper documentation was maintained for dataset preparation, preprocessing techniques, model architectures, hyperparameter selection, software implementation, testing procedures, deployment methodology, and project workflow. Comprehensive documentation ensures that the system can be understood, maintained, and extended by future developers.

Version control was maintained using Git and GitHub, providing an organized repository containing project source code, frontend implementation, backend services, trained models, and supporting documentation. GitHub facilitated systematic code management, collaboration, and version tracking throughout the development lifecycle.

The internship also resulted in the preparation of technical reports, weekly progress reports, presentation materials, and demonstration content required by both the company and the university. These documents summarize project progress, completed tasks, implementation strategies, and achieved outcomes during each phase of the internship.

Overall, the deliverables produced during this internship extend beyond a single machine learning model. They include a complete AI-powered healthcare application consisting of a prepared dataset, trained deep learning models, a responsive frontend, a robust backend, integrated prediction services, authentication mechanisms, technical documentation, deployment resources, and version-controlled source code. Together, these deliverables demonstrate the successful completion of an end-to-end software engineering project utilizing modern artificial intelligence technologies.

1.7 Report Organization

The internship involved multiple stages ranging from requirement analysis and dataset preparation to model development, software implementation, deployment, and evaluation. To present these activities in a structured and systematic manner, this report has been organized into ten chapters. Each chapter focuses on a specific phase of the project and collectively provides a comprehensive understanding of the entire development process.

Chapter 1 – Introduction presents the overall background of the project and establishes the motivation for developing an intelligent pneumonia detection system. It introduces the healthcare problem addressed during the internship, discusses the importance of artificial intelligence in medical diagnosis, defines the project objectives, explains the project scope, outlines the expected deliverables, and provides an overview of the report structure.

Chapter 2 – Requirement Analysis discusses the domain of pneumonia detection and examines the existing approaches used for diagnosing the disease. It presents a comparative study of traditional diagnostic methods and artificial intelligence-based solutions, identifies the limitations of existing systems, and explains the motivation for the proposed solution. The chapter also defines the functional and non-functional requirements of the system and introduces the software tools, programming languages, libraries, frameworks, and development environment used throughout the project.

Chapter 3 – Data Collection and Understanding describes the process of collecting the chest X-ray dataset used for model development. It explains the source of the dataset, dataset organization, image categories, directory structure, and characteristics of the collected data. The chapter also includes a detailed data dictionary and evaluates dataset quality by examining image consistency, class distribution, and potential data-related challenges.

Chapter 4 – Data Preprocessing focuses on preparing the dataset for deep learning model training. The chapter discusses image resizing, normalization, directory organization, duplicate verification, validation of dataset integrity, and preparation of training, validation, and testing datasets. Proper preprocessing ensures that the dataset is suitable for efficient model training while minimizing the possibility of errors during implementation.

Chapter 5 – Data Analysis and Modeling presents one of the most important phases of the internship. It begins with Exploratory Data Analysis (EDA), where various characteristics of the dataset are analyzed through graphical visualization and statistical observations. The chapter then explains the implementation of different deep learning architectures, including the custom CNN, MobileNetV2, and ResNet50 models. Model training procedures, hyperparameter selection, optimization techniques, evaluation metrics, and comparative performance analysis are discussed in detail. The chapter concludes with the selection of the most suitable model for deployment based on experimental results.

Chapter 6 – System Design and Development describes the software engineering aspects of the project. It explains the overall system architecture, frontend development using React, backend implementation using Flask, integration of the trained deep learning model, authentication mechanism, image upload workflow, prediction pipeline, and user interface design. This chapter illustrates how the individual software components communicate to provide a seamless user experience.

Chapter 7 – Results and Insights presents the output generated by the developed application. It includes prediction results, evaluation metrics, comparison of different models, observations obtained during experimentation, and insights derived from model performance. The chapter also discusses the practical significance of the developed system and its contribution to intelligent healthcare applications.

Chapter 8 – Testing and Validation explains the testing methodologies adopted during software development. Various testing techniques, including functional testing, integration testing, system testing, and validation of prediction accuracy, are discussed. Test cases are presented to verify different system functionalities such as authentication, image upload, prediction generation, and result visualization. The limitations of the developed system are also discussed in this chapter.

Chapter 9 – Deployment focuses on the deployment of the completed application. It describes the deployment workflow, model integration, server configuration, application execution, user accessibility, project documentation, and source code management. The chapter demonstrates how the trained deep learning model was successfully transformed into a deployable web application.

Chapter 10 – Conclusion and Future Scope summarizes the work completed during the internship and evaluates whether the defined objectives were achieved successfully. It highlights the technical knowledge acquired throughout the internship, discusses the overall contribution of the project, and suggests future enhancements such as multiclass disease classification, explainable artificial intelligence, cloud deployment, mobile application development, and integration with hospital information systems.

The report concludes with the References section, which lists all books, research papers, websites, software documentation, and online resources consulted during the project. Finally, the Appendices include supporting materials such as dataset samples, application screenshots, model summaries, hyperparameter configurations, user guide, and project timeline that complement the main content of the report.

By organizing the report in this manner, readers can systematically understand every stage of the project, from identifying the healthcare problem to designing, implementing, evaluating, and deploying an intelligent pneumonia detection system.

CHAPTER 2

REQUIREMENT ANALYSIS

2.1 Domain Study

Artificial Intelligence (AI) has become one of the most influential technologies in modern healthcare. Healthcare organizations generate enormous volumes of medical data every day through patient records, laboratory investigations, diagnostic imaging, prescriptions, wearable devices, and hospital information systems. Analyzing such large-scale data manually is time-consuming and often difficult due to the complexity of medical information. Artificial Intelligence enables computer systems to analyze healthcare data, recognize meaningful patterns, support clinical decision-making, and improve diagnostic accuracy.

Machine Learning (ML), a subset of Artificial Intelligence, allows computer systems to learn from historical data and make predictions without being explicitly programmed. Deep Learning (DL), a specialized branch of machine learning, has significantly improved image classification and pattern recognition tasks through the use of artificial neural networks with multiple hidden layers.

Medical imaging is one of the most successful applications of deep learning because medical images contain complex visual features that traditional image processing techniques often fail to capture accurately. Convolutional Neural Networks (CNNs) have revolutionized medical image analysis by automatically extracting relevant image features without requiring manual feature engineering. These capabilities have made deep learning an essential tool for disease detection and medical diagnosis.

Healthcare institutions worldwide are increasingly integrating AI-powered diagnostic systems to support physicians in disease screening, image interpretation, treatment planning, and patient monitoring. Although AI systems cannot replace medical professionals, they serve as valuable decision-support tools that improve efficiency and reduce diagnostic workload.

Overview of Pneumonia

Pneumonia is a respiratory disease caused by bacterial, viral, or fungal infections that affect the lungs. The disease primarily causes inflammation in the alveoli, the tiny air sacs responsible for oxygen exchange. During infection, these air sacs become filled with fluid or pus, reducing the lungs' ability to supply oxygen to the bloodstream.

Patients suffering from pneumonia may experience symptoms such as persistent cough, fever, chills, chest pain, rapid breathing, fatigue, shortness of breath, and reduced oxygen saturation. The severity of the disease varies depending on age, immune condition, underlying health problems, and the type of infection.

Pneumonia remains one of the leading causes of death worldwide, particularly among:

- Children below five years of age
- Elderly individuals above sixty-five years
- Patients with chronic illnesses
- Individuals with weakened immune systems
- Patients suffering from respiratory disorders

Early diagnosis significantly improves recovery rates because appropriate treatment can be initiated before severe complications develop. Therefore, rapid and accurate diagnosis remains one of the highest priorities in respiratory healthcare.

Role of Chest X-ray Imaging

Chest X-ray imaging is one of the most widely used diagnostic tools for detecting pneumonia. Compared to CT scans and MRI, chest X-rays are relatively inexpensive, quick to perform, and available in almost every healthcare facility.

Radiologists examine chest X-ray images to identify abnormalities such as:

- Lung opacity
- Pulmonary infiltrates
- Fluid accumulation
- Consolidation
- Airspace inflammation
- White patches indicating infection

These abnormalities indicate whether pneumonia may be present.

Although chest X-rays provide valuable diagnostic information, image interpretation depends heavily on the experience of radiologists. Slight abnormalities may be overlooked, especially during the early stages of infection. Furthermore, image quality, patient positioning, and exposure settings may influence diagnostic accuracy.

The increasing number of chest X-ray examinations performed daily has created a growing need for intelligent computer-aided diagnostic systems capable of assisting healthcare professionals in interpreting medical images.

Artificial Intelligence in Pneumonia Detection

Deep learning has significantly improved automatic disease detection from medical images. Instead of manually identifying image features, convolutional neural networks learn hierarchical representations directly from training images.

The workflow of AI-based pneumonia detection generally consists of:

- Collection of chest X-ray images.
- Image preprocessing.
- Data normalization.
- Model training.
- Model validation.
- Prediction generation.
- Performance evaluation.

During training, the neural network learns the visual differences between healthy lungs and lungs affected by pneumonia. After sufficient training, the model can classify new images with high accuracy.

The use of transfer learning has further accelerated medical image classification. Instead of training very large neural networks from scratch, pretrained models are fine-tuned using medical datasets. This approach reduces computational cost while improving classification accuracy.

Popular deep learning models used in pneumonia detection include:
Convolutional Neural Networks (CNN)

- MobileNetV2
- ResNet50
- DenseNet121
- EfficientNet
- InceptionV3

In this internship, three different architectures—Custom CNN, MobileNetV2, and ResNet50—were implemented and compared to identify the most suitable model for deployment.

Importance of Computer-Aided Diagnosis

Computer-Aided Diagnosis (CAD) systems are intelligent software applications that assist medical professionals in interpreting diagnostic images. These systems analyze medical images and provide predictive outputs based on trained machine learning models.

The major benefits of CAD systems include:

- Faster diagnosis
- Improved consistency
- Reduced workload for radiologists
- Early disease detection
- Increased accessibility in rural healthcare
- Support for inexperienced medical practitioners
- Improved clinical decision-making

However, CAD systems are designed to assist healthcare professionals rather than replace them. Final diagnosis should always be confirmed by qualified medical practitioners.

Need for the Proposed System

Despite significant progress in artificial intelligence research, many existing pneumonia detection models remain limited to research environments. Several published studies focus primarily on improving model accuracy but do not provide complete software solutions that healthcare professionals can easily use.

Many available systems also lack:

- User-friendly interfaces
- Secure authentication
- Image upload functionality
- Real-time prediction
- Integrated frontend and backend
- Easy deployment
- Practical usability

The proposed Pneumonia Detection System addresses these limitations by integrating deep learning with modern web technologies. The application enables users to upload chest X-ray

images through an intuitive interface, processes the images using a trained deep learning model, and returns prediction results within a few seconds.

Unlike standalone research models, the proposed system represents an end-to-end software solution incorporating machine learning, web development, backend integration, user authentication, and deployment.

Advantages of the Proposed Domain

Developing an AI-powered Pneumonia Detection System offers several technical and practical advantages.

From a healthcare perspective, the system can:

- Support radiologists during preliminary diagnosis.
- Reduce diagnosis time.
- Improve consistency in image interpretation.
- Assist hospitals experiencing heavy patient workloads.
- Provide diagnostic support in remote healthcare centers.
- From a technical perspective, the project demonstrates practical implementation of:
 - Artificial Intelligence
 - Machine Learning
 - Deep Learning
 - Computer Vision
 - Image Processing
 - React Development
 - Flask Backend Development
 - REST API Integration
 - Model Deployment
 - Full-Stack Software Engineering

The project therefore combines healthcare knowledge with modern software engineering practices to develop an intelligent decision-support system capable of assisting in pneumonia diagnosis.

2.2 Existing System Analysis

The analysis of existing systems helps identify the strengths and limitations of current pneumonia diagnosis methods and provides a foundation for developing an improved solution. This study includes both traditional diagnostic techniques and recent Artificial Intelligence (AI)-based approaches used for pneumonia detection.

2.2.1 Traditional Diagnosis

Pneumonia is traditionally diagnosed through clinical examination, laboratory tests, and chest X-ray imaging. Among these methods, chest X-rays are the most commonly used because they provide a clear view of the lungs and help identify abnormalities such as lung opacities and consolidations.

The diagnosis mainly depends on the expertise of radiologists who manually examine X-ray images. While this approach is reliable, it has several limitations. Manual interpretation is time-consuming, subject to human error, and may produce inconsistent results depending on the radiologist's experience. In hospitals with a large number of patients, delayed diagnosis may affect timely treatment. Moreover, many rural healthcare centers face a shortage of experienced radiologists, making accurate diagnosis more challenging.

2.2.2 Existing AI-Based Systems

With the advancement of Artificial Intelligence and Deep Learning, several automated pneumonia detection systems have been developed using Convolutional Neural Networks (CNNs) and transfer learning models such as MobileNetV2, ResNet50, DenseNet121, and EfficientNet. These models automatically extract image features and classify chest X-ray images with high accuracy.

Most AI-based systems follow a standard workflow that includes dataset collection, image preprocessing, model training, prediction, and performance evaluation. Although these systems achieve promising results, many are developed only for research purposes and are not integrated into complete software applications.

2.2.3 Limitations of Existing Systems

Despite their high prediction accuracy, existing AI-based systems have several limitations:

Most models are developed only for research environments such as Google Colab or Jupyter Notebook.

Many systems do not provide a user-friendly web interface.

Frontend, backend, and AI models are often not integrated into a single application.

Some models require high computational resources for deployment.

Authentication, user management, and security features are generally absent.

Most systems are difficult for non-technical users to operate.

These limitations reduce the practical usability of existing solutions in healthcare environments.

2.2.4 Comparison of Existing and Proposed System

	Traditional System	Existing System	AI Proposed System
Diagnosis Method	Manual	AI-Based	AI-Based
Chest X-ray Analysis	Manual	Automatic	Automatic
User Interface	No	Limited	Yes
Authentication	No	No	Yes
Real-Time Prediction	No	Limited	Yes
Web Application	No	Rarely	Yes
Easy to Use	Moderate	Moderate	High
Deployment Ready	No	Limited	Yes

Table 2.1
Comparison
of Existing
and
Proposed
System

2.3 Proposed System

The proposed system is an intelligent web-based application developed to assist in the early detection of pneumonia using chest X-ray images. It integrates Deep Learning techniques with modern web technologies to provide a simple, accurate, and user-friendly platform for pneumonia prediction. The system enables users to upload a chest X-ray image through a web interface, processes the image using a trained deep learning model, and displays the prediction result within a few seconds.

Unlike traditional diagnosis methods that depend entirely on manual interpretation of X-ray images, the proposed system automates the image classification process using a pre-trained deep learning model. The application is designed as a decision-support tool that helps healthcare professionals perform preliminary screening more efficiently. It is not intended to replace medical experts but to provide additional assistance during diagnosis.

The system follows a modular architecture where each component performs a specific function. The frontend, developed using React, provides an interactive interface for user registration, login, image upload, and prediction display. The backend, implemented using Flask, manages API requests, processes uploaded images, loads the trained deep learning model, and returns prediction results to the frontend.

The deep learning model is trained using a chest X-ray dataset containing two classes: Normal and Pneumonia. Before training, the dataset undergoes preprocessing steps such as image resizing, normalization, and validation to ensure consistency. Multiple models, including a Custom CNN, MobileNetV2, and ResNet50, were evaluated during development, and the best-performing model was integrated into the application for prediction.

The prediction workflow begins when a user uploads a chest X-ray image through the web application. The image is transmitted to the Flask backend, where it is resized to the required input dimensions and normalized. The preprocessed image is then passed to the trained deep learning model, which analyzes the image and predicts whether it belongs to the Normal or Pneumonia category. The prediction result is sent back to the frontend and displayed to the user in a clear and understandable format.

To improve usability and security, the application also includes user authentication. New users can register by creating an account, while existing users can securely log in to access the prediction functionality. User profile management allows authenticated users to update their personal information and maintain a personalized experience.

The proposed system has been developed with scalability in mind. Its modular architecture allows individual components, such as the frontend, backend, or AI model, to be updated independently without affecting the overall functionality. This flexibility enables future enhancements such as multiclass disease detection, cloud deployment, mobile application support, and integration with hospital information systems.

Overall, the proposed Pneumonia Detection System combines artificial intelligence, computer vision, and web development technologies to provide a complete healthcare application capable of assisting in the rapid and accurate detection of pneumonia from chest X-ray images.




Figure 2.1

2.4 Functional Requirements

Functional requirements describe the core features and operations that the system must perform to meet the project objectives. These requirements define how users interact with the application and how the system processes user requests. The proposed Pneumonia Detection System has been designed with several functional modules that work together to provide an efficient and user-friendly experience.

2.4.1 User Registration

The system allows new users to create an account by providing basic information such as name, email address, and password. During registration, the system validates the entered information to ensure correctness and uniqueness. After successful registration, user details are securely stored in the database, allowing users to access the application in future sessions.

2.4.2 User Login and Authentication

Registered users can log into the application using their email address and password. The authentication mechanism verifies the entered credentials before granting access to the system. This feature ensures that only authorized users can use the pneumonia prediction service and access their personal profiles.

2.4.3 Image Upload

The application provides an interface through which authenticated users can upload chest X-ray images. The system accepts images in supported formats such as JPG, JPEG, and PNG. Once uploaded, the image is transferred securely to the backend server for processing.

2.4.4 Image Preprocessing

Before prediction, the uploaded image is automatically preprocessed by the backend. This process includes resizing the image to the required input dimensions, converting it into the appropriate format, and normalizing pixel values. These preprocessing steps ensure compatibility with the trained deep learning model and improve prediction accuracy.

2.4.5 Pneumonia Prediction

After preprocessing, the image is passed to the trained deep learning model. The model analyzes the chest X-ray and classifies it into one of two categories:

- Normal
- Pneumonia

The prediction result is then returned to the backend for further processing.

2.4.6 Display of Prediction Result

The backend sends the prediction result to the React frontend, where it is displayed in a clear and user-friendly manner. The application informs the user whether the uploaded X-ray image is predicted as Normal or Pneumonia, enabling quick preliminary assessment.

2.4.7 Profile Management

The system allows authenticated users to manage their personal information. Users can update profile details such as their name, email address, and password whenever required. This feature enhances user convenience and maintains accurate account information.

2.4.8 Logout

The application provides a logout option that securely ends the user's session. This prevents unauthorized access and improves overall system security, especially when the application is used on shared devices.

2.5 Non-Functional Requirements

Non-functional requirements define the quality attributes of the system. They ensure that the application performs efficiently, remains secure, and provides a satisfactory user experience.

2.5.1 Performance

The system should generate prediction results within a few seconds after an image is uploaded. Efficient preprocessing and model inference help minimize response time and provide a smooth user experience.

2.5.2 Reliability

The application should consistently produce prediction results without unexpected failures. Proper exception handling and validation mechanisms improve system reliability during operation.

2.5.3 Security

Since the application includes user authentication, user credentials and personal information must be stored securely. Only authenticated users should be allowed to access prediction services. Secure session management further protects user accounts.

2.5.4 Usability

The user interface should be simple, responsive, and easy to navigate. Even users with limited technical knowledge should be able to upload images and obtain prediction results without difficulty.

2.5.5 Scalability

The modular architecture of the application allows future enhancements such as multiclass disease classification, cloud deployment, mobile application support, and integration with hospital management systems.

2.5.6 Maintainability

The frontend, backend, and deep learning model are developed as independent modules. This modular design simplifies maintenance, debugging, and future upgrades without affecting the overall system.

2.5.7 Compatibility

The application should function correctly on modern web browsers such as Google Chrome, Microsoft Edge, Mozilla Firefox, and Brave. It should also support different screen sizes for desktop and laptop users.




Figure 2.2

2.6 Tools and Technologies Used

The development of these complex, full-stack AI systems required the mastery of a diverse, modern technology stack:

Tool / Technology	Purpose
Tool / Technology	Purpose
Python	Backend programming and AI development
TensorFlow	Deep learning framework
Keras	Neural network implementation
MobileNetV2	Transfer learning model
ResNet50	Model comparison
OpenCV	Image preprocessing
NumPy	Numerical computations
Pandas	Data handling and preprocessing
React	Frontend development
Flask	Backend API development
Google Colab	Model training
VSCode	Code development
Git	Version control
GitHub	Source code management

CHAPTER 3

DATA COLLECTION AND UNDERSTANDING

3.1 Data Source

The success of any Artificial Intelligence model depends largely on the quality of the training dataset. For this project, a publicly available Chest X-ray Images (Pneumonia) dataset was selected because it contains labeled X-ray images suitable for binary image classification tasks.

The dataset was obtained from Kaggle, one of the most widely used platforms for machine learning and data science projects. It contains chest X-ray images collected from pediatric patients and is commonly used by researchers to evaluate deep learning models for pneumonia detection.

The dataset is organized into three main folders:

- Training Dataset (Train) – Used to train the deep learning models.
- Validation Dataset (Val) – Used during model training to evaluate learning performance and prevent overfitting.
- Testing Dataset (Test) – Used after training to evaluate the final performance of the model on unseen data.

Each of these folders contains two categories of images:

- Normal
- Pneumonia

This directory structure allows deep learning frameworks such as TensorFlow and Keras to automatically identify image classes during training.

The dataset was selected because it provides a sufficient number of labeled medical images, making it suitable for developing and evaluating deep learning models such as CNN, MobileNetV2, and ResNet50.

Dataset Details

Attribute	Description
Dataset Name	Chest X-ray Images (Pneumonia)
Source	Kaggle
Data Type	Medical Images
Image Format	JPEG (.jpeg)
Number of Classes	2
Categories	Normal, Pneumonia
Image Type	Chest X-ray
Application	Binary Image Classification

3.2 Data Description

The collected dataset contains chest X-ray images representing two different medical conditions: healthy lungs and lungs affected by pneumonia. Each image is labeled according to its corresponding category, making the dataset suitable for supervised learning.

The Normal class consists of chest X-ray images that show healthy lungs without signs of infection or abnormal opacity. These images serve as the reference class during model training.

The Pneumonia class contains chest X-ray images showing visible signs of lung infection. These images typically exhibit features such as lung opacity, consolidation, or inflammation, which are commonly associated with pneumonia. The deep learning model learns to distinguish these patterns from healthy lungs during the training process.

The dataset contains images of varying resolutions because they were collected from different medical sources. Since deep learning models require a fixed input size, all images were resized during preprocessing to ensure consistency before training.

Another important characteristic of the dataset is that it includes images with different brightness levels, contrast variations, and minor positional differences. These natural variations improve the model's ability to generalize and make accurate predictions on previously unseen chest X-ray images.

Overall, the dataset provides sufficient diversity to train a robust deep learning model while maintaining a clear binary classification structure.

3.3 Data Structure

The dataset follows a hierarchical folder structure that separates images based on their purpose and class labels. Organizing the dataset in this format simplifies the loading process using TensorFlow's `image_dataset_from_directory()` function and improves the efficiency of model training.

The overall directory structure is shown below.

```
Chest_Xray_Dataset/  
├── train/  
│   ├── NORMAL/  
│   └── PNEUMONIA/  
├── val/  
│   ├── NORMAL/  
│   └── PNEUMONIA/  
└── test/  
    ├── NORMAL/  
    └── PNEUMONIA/
```

This explicit directory structure is critical. It allows the `ImageDataGenerator` class in Keras to utilize the `flow_from_directory` method. This method automatically infers the binary class labels based strictly on the subfolder names, entirely eliminating the need for complex, manual CSV label mapping and significantly streamlining the data loading process.

3.4 Data Dictionary (MediaPipe Spatial Outputs)

A data dictionary provides information about the different attributes and components of the dataset. Since this project involves image data rather than tabular data, the dictionary describes the characteristics of each image and its associated label.

Attribute	Description
Image ID	Unique identifier of the image file
Image Name	Name of the X-ray image
Image Format	JPEG (.jpeg)
Image Size	Variable (resized during preprocessing)
Image Type	Chest X-ray
Class Label	Normal / Pneumonia
Dataset Category	Train / Validation / Test

Table 3.2 Data Dictionary

The dataset does not contain additional patient information such as age, gender, or medical history. It is specifically designed for image classification based on visual features present in chest X-ray images.

3.5 Data Quality Assessment

Before training the deep learning models, the dataset was carefully examined to assess its quality and identify any issues that might affect model performance. Data quality assessment is an important step because poor-quality data can lead to inaccurate predictions and reduced model reliability.

The assessment focused on several aspects of the dataset.

Image Quality

The chest X-ray images were visually inspected to ensure that they were clear and suitable for analysis. Most images were of good quality and contained sufficient detail for identifying lung abnormalities.

Class Labels

Each image was verified to ensure that it was correctly organized into either the Normal or Pneumonia category. Proper labeling is essential for supervised learning because incorrect labels can negatively impact model accuracy.

Missing Data

The dataset was checked for missing or corrupted image files. No significant missing data was observed, allowing the dataset to be used without major modifications.

Duplicate Images

The dataset was reviewed for duplicate image files that could introduce bias during model training. Any unnecessary duplicate images were excluded where applicable to improve model generalization.

Image Resolution

Since the images were collected from multiple sources, they had different resolutions. To maintain consistency, all images were resized to the required input dimensions during preprocessing before being fed into the deep learning models.

Dataset Consistency

The folder organization, file naming, and image formats were verified to ensure compatibility with TensorFlow's dataset loading functions. The dataset maintained a consistent directory structure, making it suitable for automated loading and preprocessing.

After completing the quality assessment, the dataset was considered suitable for training deep learning models. Minor variations in image size and brightness were handled during preprocessing, ensuring that the final prepared dataset met the requirements of the proposed Pneumonia Detection System.

CHAPTER 4

DATA PREPROCESSING

4.1 Data Cleaning

Data cleaning is the process of identifying and correcting inconsistencies in the dataset before model training. Since this project uses image data instead of structured tabular data, the cleaning process focused on ensuring that all chest X-ray images were suitable for analysis.

Initially, the dataset folders were examined to verify that images were correctly organized into the Normal and Pneumonia classes. The directory structure was checked to ensure compatibility with TensorFlow's image loading functions.

Each image was also inspected to confirm that it could be successfully loaded without errors. Corrupted or unreadable image files, if any, were removed to prevent interruptions during training. Images with unsupported formats were excluded, ensuring that only valid image files were included in the dataset.

The file naming conventions and folder organization were also reviewed to maintain consistency across the training, validation, and testing datasets. These cleaning activities ensured that the dataset was reliable and ready for further preprocessing.

4.2 Handling Missing Values

Unlike tabular datasets, image datasets generally do not contain missing values in the form of empty cells or null entries. However, missing data may occur when image files are absent, damaged, or unreadable.

During preprocessing, the dataset was verified to ensure that all expected image files were present in their respective directories. Images that could not be opened or were corrupted were identified and excluded from the dataset. Since the selected chest X-ray dataset was well-maintained, no significant missing image files were found.

The class labels were also verified to ensure that each image belonged to the correct category. This validation helped maintain data integrity throughout the training process.

By confirming the completeness of the dataset, the risk of training interruptions and incorrect predictions was minimized.

4.3 Duplicate Removal

Duplicate images can introduce bias into a deep learning model by causing the model to learn the same visual patterns repeatedly. This may lead to overfitting, where the model performs well on training data but poorly on unseen images.

The dataset was reviewed to identify duplicate image files. The folder structure and image filenames were examined to ensure that unnecessary duplicate images were not present across the training, validation, and testing datasets.

Maintaining unique images in each dataset improves the model's ability to generalize and reduces the likelihood of overly optimistic performance during evaluation.

Although the selected dataset contained very few duplicate images, this verification step was important to maintain the quality and reliability of the training data.

4.4 Data Transformation

Data transformation converts raw images into a format suitable for deep learning models. Since convolutional neural networks require standardized input dimensions, all chest X-ray images were transformed before training.

The following preprocessing operations were performed:

Image Resizing

The original chest X-ray images had different resolutions because they were collected from multiple medical sources. To ensure consistency, every image was resized to 224×224 pixels, which matches the input size required by MobileNetV2 and ResNet50.

Pixel Normalization

Pixel values in digital images typically range from 0 to 255. These values were normalized to the range 0 to 1 by dividing each pixel value by 255. Normalization improves training stability and allows the neural network to converge more efficiently.

Image Format Conversion

All uploaded images were converted into RGB format before being passed to the deep learning model. This ensured compatibility with the pretrained architectures used in the project.

Dataset Splitting

The dataset was organized into separate training, validation, and testing sets. This division ensured that the model was trained, validated, and evaluated on different subsets of data, allowing an unbiased assessment of model performance.

These transformation steps ensured that every image followed a consistent format before entering the model training pipeline.

4.5 Data Validation & Augmentation

After preprocessing, the prepared dataset was validated to confirm that all images satisfied the model requirements.

The validation process included the following checks:

- Verification of the folder structure.
- Confirmation that every image belonged to the correct class.
- Validation of image dimensions after resizing.
- Checking successful loading of all images.
- Verification of normalized pixel values.
- Ensuring compatibility with TensorFlow data loading functions.
-

A sample of preprocessed images was also visually inspected to verify that resizing and normalization did not distort important medical features. This quality assurance step helped ensure that clinically relevant information was preserved while preparing the dataset.

Successful validation confirmed that the dataset was ready for model training.

4.6 Final Prepared Dataset


After completing all preprocessing activities, the final dataset was prepared for training the deep learning models. The prepared dataset contained standardized chest X-ray images organized into training, validation, and testing directories.

Each image had undergone:

- Data cleaning
- Image verification
- Duplicate checking
- Image resizing
- Pixel normalization
- Dataset validation

The final dataset served as the input for training and evaluating the Custom CNN, MobileNetV2, and ResNet50 models.

Preparing a high-quality dataset significantly improved the reliability of the developed Pneumonia Detection System and contributed to better prediction accuracy during model evaluation and deployment.



CHAPTER 5

DATA ANALYSIS AND MODELING

5.1 Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an important step that helps understand the dataset before model training. It involves examining the structure, distribution, and quality of the data to identify patterns or inconsistencies that may influence model performance.

Since this project uses image data instead of numerical datasets, EDA focused on analyzing image characteristics, dataset organization, and class distribution.

Class Distribution

The first step was to examine the number of images available in each class. The dataset consists of two categories:

- Normal
- Pneumonia

Understanding the distribution of images helps determine whether the dataset is balanced or imbalanced. A balanced dataset generally allows the model to learn both classes effectively, whereas a highly imbalanced dataset may introduce prediction bias.

A graphical representation of the class distribution was created to visualize the number of images in each category.

The analysis showed that both categories contained sufficient images for training the deep learning models, making the dataset suitable for binary image classification.

Sample Image Analysis

Several chest X-ray images from both classes were visually inspected before training.

The Normal class images generally displayed clear lung fields without significant abnormalities. In contrast, the Pneumonia class images showed visible lung opacities, consolidation, and cloudy regions caused by infection.

This visual inspection confirmed that meaningful differences existed between the two categories, allowing deep learning models to learn discriminative features during training.

Image Resolution Analysis

The original dataset contained images of different resolutions because they were collected from multiple healthcare sources.

Since deep learning models require uniform input dimensions, image sizes were analyzed before preprocessing. To maintain consistency, every image was resized to 224×224 pixels, ensuring compatibility with MobileNetV2 and ResNet50 architectures.

This preprocessing step reduced computational complexity while preserving essential medical features.

Data Quality Verification

EDA also included verification of image quality and dataset consistency.

The following checks were performed:

- Image readability
- Folder organization
- Label consistency
- Image format verification
- Duplicate image inspection

The dataset was found to be suitable for model development after preprocessing.

5.2 Data Relationships (Spatial Kinematics)

Unlike structured datasets containing numerical attributes, image datasets represent relationships through visual patterns rather than statistical correlations.

In chest X-ray images, the relationship between the input image and the target class depends on several visual characteristics, including:

- Lung opacity
- Texture variations
- Image contrast
- Brightness distribution
- Shape of infected regions

Deep learning models automatically identify these hidden relationships during training through convolutional operations.

Instead of manually extracting image features, convolutional neural networks learn hierarchical representations of medical images.

The learning process generally follows multiple stages:

- Detection of edges
- Recognition of textures
- Identification of lung structures
- Detection of abnormal regions
- Final classification

This automatic feature extraction significantly improves classification accuracy compared to traditional image processing methods.

5.3 Key Metrics Identification

The performance of deep learning models cannot be evaluated using accuracy alone. Therefore, multiple evaluation metrics were used to assess the effectiveness of the proposed system.

Accuracy

Accuracy represents the percentage of correctly classified chest X-ray images.

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN}$$

where:


TP = True Positive

TN = True Negative

FP = False Positive

FN = False Negative

Accuracy provides an overall measure of model performance.



Precision

Precision measures how many images predicted as pneumonia are actually pneumonia.

$$Precision = \frac{TP}{TP+FP}$$

High precision indicates a lower number of false positive predictions.

Recall

Recall measures the ability of the model to correctly identify pneumonia cases.

$$Recall = \frac{TP}{TP+FN}$$

High recall is particularly important in medical diagnosis because missing a pneumonia case may delay treatment.

F1-Score

The F1-score combines Precision and Recall into a single evaluation metric.

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

A higher F1-score indicates balanced predictive performance.

Confusion Matrix

The Confusion Matrix provides a detailed summary of classification results by comparing predicted labels with actual labels.

It consists of:

- True Positive
- True Negative
- False Positive
- False Negative

The confusion matrix helps identify the types of prediction errors made by the model.

5.4 Business Rules and Logic (Geometric Heuristics)

The Pneumonia Detection System follows a predefined workflow that ensures every uploaded image is processed consistently before prediction.

The system logic is described below:

-
- User logs into the application.
 - User uploads a chest X-ray image.
 - Backend validates the uploaded file.
 - Image is resized to 224×224 pixels.
 - Pixel values are normalized.
 - Preprocessed image is passed to the trained MobileNetV2 model.
 - Model generates prediction probabilities.
 - Highest probability determines the final class.
 - Result is displayed as Normal or Pneumonia.
 - User receives prediction through the web interface.
-

This workflow ensures efficient and reliable prediction while maintaining consistency across all uploaded images.

5.5 Data Modeling Approach (The Scaling Experiment)

The primary objective of the modeling phase was to develop a deep learning model capable of accurately classifying chest X-ray images.

Three different architectures were implemented and evaluated during the internship.

Custom Convolutional Neural Network (CNN)

The first model developed was a Custom CNN consisting of convolutional layers, pooling layers, activation functions, and fully connected layers.

The model automatically learned visual features from chest X-ray images through multiple convolutional operations.

Although the CNN demonstrated satisfactory performance, its prediction accuracy was lower than the transfer learning models.

MobileNetV2

MobileNetV2 was selected as the primary transfer learning model because of its lightweight architecture and high computational efficiency.

A pretrained MobileNetV2 model was loaded with ImageNet weights, while the final classification layer was modified for binary classification.

Initially, the pretrained layers were frozen, allowing only the newly added classification layers to be trained.

After initial training, fine-tuning was performed by unfreezing the last few layers of the network. This enabled the model to learn pneumonia-specific features while retaining the general visual knowledge acquired during pretraining.

The model demonstrated excellent performance in terms of accuracy, convergence speed, and computational efficiency, making it the preferred choice for deployment.

ResNet50

ResNet50 was implemented as another transfer learning model for comparison.

Its residual learning architecture enables very deep neural networks to train effectively without suffering from the vanishing gradient problem.

Although ResNet50 achieved competitive performance, it required higher computational resources and longer training time compared to MobileNetV2.

Therefore, MobileNetV2 was selected for deployment because it provided a better balance between prediction accuracy and inference speed.

Model Name	Dataset Size	Training Time (sec)	Accuracy	Loss
CNN	200	249.91 seconds	0.9	0.4287
CNN	500	1416.64 seconds	0.56	0.6928
CNN	700	979.98 seconds	0.9	0.4147
CNN	1000	1334.39 seconds	0.975	0.0633
Model Name	Dataset Size	Training Time (sec)	Accuracy	Loss
MobileNetV2	200	67.92 seconds	0.875	0.3253
MobileNetV2	500	146.81 seconds	0.96	0.1746
MobileNetV2	700	203.28 seconds	0.95	0.1537
MobileNetV2	1000	362.89 seconds	0.94	0.1212
Model Name	Dataset Size	Training Time (sec)	Accuracy	Loss
ResNet50	200	205.27 seconds	0.675	0.6699
ResNet50	500	431.91 seconds	0.45	0.6661
ResNet50	700	716.74 seconds	0.7786	0.609
ResNet50	1000	898.55 seconds	0.795	0.5791

Figure:5.2

CHAPTER 6

DASHBOARD DESIGN AND DEVELOPMENT

6.1 Dashboard Planning

Before developing the application, the overall system workflow and user requirements were analyzed. The planning phase focused on designing a web application that would be simple to use while ensuring efficient communication between the frontend, backend, and deep learning model.

The application was planned with the following objectives:

- Provide a clean and responsive user interface.
- Allow secure user registration and login.
- Enable users to upload chest X-ray images easily.
- Process uploaded images efficiently using the trained model.
- Display prediction results clearly and quickly.
- Maintain a modular structure for future enhancements.

The system was divided into independent modules, allowing each component to perform a specific task without affecting other parts of the application.

The major modules include:

- User Authentication Module
- Image Upload Module
- Image Preprocessing Module
- Prediction Module
- Result Display Module
- Profile Management Module

This modular design improves maintainability and simplifies future development.

6.2 Data Model Design (API Specifications)

The application follows a three-layer architecture consisting of the presentation layer, application layer, and intelligence layer.

Presentation Layer

The presentation layer is developed using React. It provides the graphical user interface through which users interact with the application. The frontend includes pages for registration, login, profile management, image upload, and prediction results.

Application Layer

The application layer is implemented using Flask. It acts as the communication bridge between the frontend and the deep learning model. It receives uploaded images, preprocesses them, performs model inference, and sends prediction results back to the frontend.

Intelligence Layer

The intelligence layer consists of the trained MobileNetV2 deep learning model. This model analyzes chest X-ray images and predicts whether the uploaded image belongs to the Normal or Pneumonia category.

The interaction between these three layers enables efficient processing while keeping the system organized and scalable.


Akshar AI
Intelligent Healthcare Solutions

Secure Gateway

PATIENT

About Akshar AI

Intelligent Healthcare Automation Solutions.

Our Mission

We aim to empower healthcare providers and patients by offering an accessible, AI-powered tool for early pneumonia detection from chest X-rays.

Core Solutions

Instant Pneumonia Detection

Fast, automated preliminary screening.

Secure Image Processing

Your medical data remains private.

Diagnostic Assistance

Aiding medical professionals with AI insights.

Clinical Disclaimer

System functions as auxiliary decision support. Findings must be validated by a licensed physician.

Patient Sign In

Sign in to access your health automation tools.

Sign In

Register

Email ID

patient@akshar.ai

Password

Forgot password?

Sign In to Portal

Don't have an account? Register here

Secure Patient Connection • Encrypted Storage


Akshar AI | Intelligent Healthcare Solutions

Developed by Yashas Raj R
Subject to Akshar AI Programme Terms & Conditions


Akshar AI
Intelligent Healthcare Solutions

Home
About
uploads
Records
Profile

Yashas Raj R

Clinical Radiography Records

Review, filter, and audit all convolutional neural network diagnostic scan histories captured inside the active session context.


TOTAL ANALYZED
3

HEALTHY CLEARANCES
1

INFILTRATIONS FLAGGED
2

Search by name, Patient ID or result...
Clear All History

File Name	Patient ID	Scan Date	AI Result	Confidence	Model	Processing Time	Report	Actions
IM-0075-0...	PNE-2026-000001	2026-07-12T14:40:31.184949	Normal	38.2%	MobileNet V2	2.40s		
person24_v...	EMP-2026-00003	2026-07-12T14:37:48.593644	Pneumonia Detected	99.9%	MobileNet V2	2.40s		
PNE-2026-	2026-07-		Pneumonia		MobileNet			



6.3:Dashboard Components

The application consists of several components that work together to provide the complete prediction workflow.

Login Module

The login module authenticates registered users before allowing access to the application. Users enter their email address and password, which are verified by the backend.

Registration Module

New users can create an account by providing their basic details. The system validates the entered information before storing it securely.

Home Page

The home page introduces the application and provides navigation to various sections. It acts as the starting point for users after logging into the system.

Image Upload Module

This module allows users to upload chest X-ray images from their local device. The uploaded image is validated before being sent to the backend for prediction.

Supported formats include:

- JPG
- JPEG
- PNG

Prediction Module

After receiving the uploaded image, the backend performs preprocessing and forwards the image to the trained MobileNetV2 model.

The model predicts one of the following classes:

- Normal
- Pneumonia


The prediction is returned to the frontend within a few seconds.

Result Display Module

The prediction result is displayed in an easy-to-understand format. Users can clearly identify whether the uploaded X-ray image indicates pneumonia or appears normal.

Profile Module

Authenticated users can view and update their personal information through the profile page.



6.4 Visualizations Used

Although this project is not a traditional business intelligence dashboard, several visual elements were incorporated to improve user interaction and understanding.

The application includes:

- Home page with intuitive navigation.
- Login and registration forms.
- Image preview before prediction.
- Prediction result screen.
- Responsive layouts for different screen sizes.
- Icons and status indicators for better user experience.

Additionally, training graphs generated during model development were used to evaluate the learning performance of the deep learning model.

These include:

-
- Training Accuracy Graph
 - Validation Accuracy Graph
 - Training Loss Graph
 - Validation Loss Graph
 - Confusion Matrix
-

These visualizations helped monitor model performance and identify potential overfitting or underfitting during training.

6.5 Interactive Features (Asynchronous JavaScript)

The Pneumonia Detection System provides several interactive features that improve usability and overall user experience.

The key interactive features include:

-
- Secure user registration and login.
 - Image upload using the file selection interface.
 - Image preview before submission.
 - Automatic preprocessing of uploaded images.
 - Real-time prediction generation.
 - Display of prediction results.
 - User profile management.
 - Responsive interface compatible with different screen sizes.
-

The application communicates with the backend using API requests, ensuring smooth data transfer between the frontend and the deep learning model.

6.6 User Interface Design (Responsive CSS)

The user interface was designed with simplicity, accessibility, and responsiveness as the primary objectives. Since the application may be used by individuals with varying levels of technical expertise, the interface emphasizes ease of use and minimal user interaction.

The design follows modern web development principles by maintaining consistent layouts, intuitive navigation, and responsive components across all pages.

The key UI screens include:

Home Page

The home page introduces the application and provides navigation to the main features.

Login Page

Allows registered users to securely access the application.

Registration Page

Enables new users to create an account.

Upload Page

Allows users to upload chest X-ray images for analysis.

Prediction Page

Displays the uploaded image along with the prediction result generated by the deep learning model.

Profile Page

Allows users to view and update their account information.

The interface was developed using reusable React components, ensuring consistency across different pages while simplifying future maintenance and updates.

CHAPTER 7

RESULTS AND INSIGHTS

7.1 Dashboard Output

The developed web application successfully integrates the trained deep learning model with the user interface, enabling users to perform pneumonia prediction through a simple workflow. After logging into the application, users can upload a chest X-ray image using the image upload page. Once the image is submitted, it is sent to the Flask backend, where preprocessing and model inference are performed. The prediction result is then displayed on the frontend.

The application provides a clean and responsive interface, allowing users to complete the prediction process without requiring technical knowledge. The result screen clearly indicates whether the uploaded image is classified as Normal or Pneumonia, enabling quick interpretation.

The prediction workflow performed by the application is summarized below:

- User logs into the application.
- Chest X-ray image is uploaded.
- Image is validated and preprocessed.
- MobileNetV2 model performs prediction.
- Prediction result is displayed on the screen.

The successful integration of all system components demonstrates that the application functions as intended and provides real-time prediction capability.




Figure 7.1: Prediction Result Screen

7.2 Key Findings (Solving the Overfitting Crisis)

Several important observations were made during model development and evaluation.

The preprocessing stage significantly influenced the overall performance of the deep learning models. Standardizing image size and normalizing pixel values helped the models learn more effectively and improved convergence during training.

Among the implemented models, MobileNetV2 provided the best balance between classification accuracy, computational efficiency, and prediction speed. Compared to the custom CNN,

MobileNetV2 achieved better performance due to its pretrained feature extraction capability. Although ResNet50 also produced good results, it required more computational resources and longer training time.

Another important finding was that transfer learning considerably reduced training time while maintaining high prediction accuracy. Instead of training a large neural network from scratch, using pretrained weights allowed the model to learn disease-specific features more efficiently.

The web application also performed reliably during testing. The integration between React, Flask, and the trained model enabled smooth communication and real-time prediction without significant delays.

Overall, the project demonstrated that deep learning models can effectively assist in the classification of chest X-ray images and can be successfully deployed within a user-friendly web application.




Figure 7.2: Sample Prediction (Normal)




Figure 7.3: Sample Prediction (Pneumonia)

7.3 Trend Analysis (The Power of Transfer Learning)

During model training, the performance of the implemented models was monitored using training and validation metrics. The accuracy and loss graphs provided valuable insights into the learning behavior of the models.

Initially, the training accuracy increased rapidly as the models learned basic image features. As training progressed, both training and validation accuracy gradually stabilized, indicating that the models were learning meaningful representations of the chest X-ray images.

Similarly, the training and validation loss decreased over successive epochs, showing that the models were improving their predictions while minimizing classification errors.

Fine-tuning the MobileNetV2 model further improved performance by allowing the last few layers of the pretrained network to adapt to the chest X-ray dataset. This resulted in better generalization and more accurate predictions on unseen images.

The observed trends confirmed that the preprocessing techniques, transfer learning approach, and model training strategy were effective in developing a reliable pneumonia detection model.

7.4 Comparative Analysis (MobileNetV2 vs. ResNet50)

To identify the most suitable model for deployment, three different deep learning architectures were implemented and compared based on prediction performance, computational efficiency, and inference speed.

The *Custom CNN* served as the baseline model. It successfully learned the basic features of chest X-ray images but produced comparatively lower classification accuracy than the transfer learning models.

ResNet50 demonstrated strong predictive capability due to its deep residual architecture. However, the model required greater computational resources and longer training time, making it less suitable for lightweight deployment.

MobileNetV2 achieved the best overall performance. Its lightweight architecture, combined with transfer learning, enabled fast training, efficient inference, and high prediction accuracy. These characteristics made it the preferred model for integration into the web application.

Table 7.1 Model Comparison

Model	Performance	Training Time	Computational Cost	Deployment Suitability
Custom CNN	Good	Moderate	Low	Moderate
MobileNetV2	Excellent	Low	Low	High
ResNet50	Very Good	High	High	Moderate

Based on this comparison, MobileNetV2 was selected as the final model for deployment because it provided the best combination of prediction accuracy, computational efficiency, and response time.

7.5 Generated Insights (Computer Vision Realities)

The implementation of the Pneumonia Detection System provided several valuable technical and practical insights.

The project demonstrated that high-quality data preprocessing plays a significant role in improving model performance. Standardizing image dimensions, validating image quality, and normalizing pixel values contributed to better learning and more stable predictions.

The use of transfer learning proved to be highly effective for medical image classification. Pretrained models such as MobileNetV2 significantly reduced the amount of training time while maintaining strong predictive performance.

The integration of the trained model into a web application showed that artificial intelligence can move beyond research environments and become accessible through practical software solutions. The combination of React and Flask provided an efficient architecture for connecting users with the deep learning model.

Another important insight was the importance of modular software design. Separating the frontend, backend, and AI model into independent components improved maintainability and simplified future enhancements.

Finally, the project highlighted the growing role of Artificial Intelligence in healthcare. Although the developed system is intended as a decision-support tool rather than a replacement for medical professionals, it demonstrates how intelligent software can assist in faster and more consistent preliminary diagnosis.

CHAPTER 8

TESTING AND VALIDATION

8.1 Testing Strategy

A systematic testing strategy was adopted to evaluate the functionality and reliability of the developed system. Testing was carried out in different stages to ensure that both the software components and the AI model operated correctly.

The following testing approaches were used:

- Functional Testing
- Integration Testing
- Model Validation
- User Interface Testing
- Performance Testing

Each module of the application was tested independently before performing complete system testing. This approach helped identify and resolve issues early in the development process.

8.2 Test Cases

Table 8.1 Functional Test Cases

Test Case ID	Module	Test Description	Expected Result	Status
TC-01	Registration	Register a new user	Account created successfully	Pass
TC-02	Login	Login with valid credentials	User logged in successfully	Pass
TC-03	Login	Login with invalid credentials	Error message displayed	Pass
TC-04	Image Upload	Upload valid chest X-ray image	Image uploaded successfully	Pass
TC-05	Prediction	Predict uploaded image	Prediction displayed	Pass
TC-06	Profile	Update profile details	Information updated	Pass
TC-07	Logout	Logout from system	Session ended successfully	Pass

The successful execution of all functional test cases confirms that the application meets its functional requirements.

8.3 Validation Process

Model validation was performed to evaluate the predictive capability of the trained deep learning models. The testing dataset, which was not used during training, served as the basis for evaluating model performance.

The validation process involved:

- Loading unseen chest X-ray images.
- Preprocessing each image.
- Passing the image to the trained model.
- Comparing predicted labels with actual labels.
- Calculating evaluation metrics.

The MobileNetV2 model demonstrated consistent prediction performance across the testing dataset and produced reliable classification results.

The evaluation confirmed that the model generalized well to unseen data, indicating effective learning during training.

8.4 Performance Evaluation

The trained models were evaluated using standard classification metrics commonly used in machine learning and medical image analysis.

The following metrics were considered:

Accuracy

Measures the percentage of correctly classified images.

Precision

Measures the proportion of correctly predicted pneumonia cases among all images predicted as pneumonia.

Recall

Measures the ability of the model to identify actual pneumonia cases.

F1-Score

Represents the harmonic mean of Precision and Recall, providing a balanced evaluation of model performance.

Confusion Matrix

Provides a detailed breakdown of correct and incorrect predictions, allowing analysis of classification errors.

These evaluation metrics provided a comprehensive understanding of model performance and confirmed the effectiveness of the selected MobileNetV2 architecture.

8.5 Limitations

Although the developed Pneumonia Detection System performed well during testing, several limitations were identified.

The system supports only binary classification (Normal and Pneumonia).

Prediction accuracy depends on the quality of the uploaded chest X-ray image.

The model was trained using a publicly available dataset, which may not represent all real-world clinical scenarios.

The application is intended as a decision-support tool and should not replace professional medical diagnosis.

Additional training using larger and more diverse datasets could further improve model performance.

Recognizing these limitations provides opportunities for future improvements and enhancements.

CHAPTER 9

DEPLOYMENT

9.1 Deployment Process

The deployment process involved transitioning the codebase from an exploratory state into a frozen, production-ready state.

The critical first step was dependency locking. Utilizing the `pip freeze > requirements.txt` command, the exact versions of every Python library used during development (e.g., TensorFlow 2.14, Flask 3.0, OpenCV 4.8) were locked. This guarantees that the deployment environment will perfectly mirror the development environment, protecting the software from breaking changes caused by future library updates.

The Flask application was structured and configured to run robustly, laying the architectural groundwork to be easily migrated behind a production-grade WSGI server (like Gunicorn or uWSGI) and a reverse proxy (like Nginx) for enterprise-scale deployment.

9.2 Publishing and Sharing

The final publishing phase involved a comprehensive cleanup of the Git repository. All deprecated Jupyter notebooks, experimental scripts, and redundant data files from earlier weeks were purged.

The pristine codebase—including the frontend HTML/CSS/JS, the backend Flask API (`app.py`), the optimal `.h5` model weights, and the computer vision tracking scripts—was committed and pushed to the master branch of the organization's central repository, officially marking the handover of the intellectual property.

9.3 User Access and Usage

Accessing the system is designed to be frictionless. For the medical dashboard, users (doctors or technicians) simply require a standard, modern web browser. They navigate to the deployed URL (or `localhost:5000` during local execution). The intuitive drag-and-drop interface requires zero technical training to operate.

For developers or engineers wishing to run the Hand Tracking module, access is granted via the command line (`python hand_tracker.py`), which immediately initializes the webcam and launches the OpenCV rendering window.

9.4 Documentation

Comprehensive, developer-empathetic documentation is the cornerstone of professional deployment. Without it, software quickly becomes unmaintainable legacy code.

A master `README.md` was authored for the repository. It explicitly details the prerequisites, the exact terminal commands required to initialize a Python virtual environment, and the steps to install the locked dependencies. Crucially, it provides a highly detailed API Specification for the `/predict` endpoint, documenting the required HTTP methods, the multipart/form-data payload structure, and providing exact examples of both successful and error JSON responses. This ensures the Akshar AI team can easily inherit, maintain, and scale the systems.

CHAPTER 10

CONCLUSION AND

FUTURE SCOPE

10.1 Project Summary

The Pneumonia Detection System developed during this internship demonstrates the successful integration of deep learning techniques with modern web technologies to create an intelligent healthcare application. The project aimed to develop an automated system capable of detecting pneumonia from chest X-ray images using deep learning models while providing a user-friendly web interface for easy accessibility.

Throughout the project, various stages of the software development lifecycle were successfully completed, including requirement analysis, data collection, preprocessing, exploratory data analysis, model development, testing, deployment, and evaluation. The implementation of multiple deep learning architectures enabled a comparative study, leading to the selection of MobileNetV2 as the final model due to its high prediction accuracy, computational efficiency, and suitability for real-time deployment.

The developed web application allows authenticated users to upload chest X-ray images and obtain prediction results through an intuitive interface. The integration of the React frontend, Flask backend, and MobileNetV2 model demonstrates how artificial intelligence can be effectively incorporated into practical software systems to assist in preliminary medical diagnosis.

Although the system is not intended to replace professional medical expertise, it serves as a valuable decision-support tool that can assist healthcare professionals by providing faster preliminary analysis of chest X-ray images. The project also provided practical experience in deep learning, web application development, image processing, software engineering, and deployment technologies, contributing significantly to both technical knowledge and professional development.

Overall, the internship successfully achieved its objectives by delivering a functional, reliable, and scalable Pneumonia Detection System that highlights the potential of artificial intelligence in improving healthcare services.

10.2 Project Achievements

The major achievements of the project are summarized below:

- Successfully studied the application of Artificial Intelligence in medical image analysis.
- Collected and analyzed a publicly available chest X-ray dataset.
- Performed comprehensive data preprocessing and validation.
- Implemented and compared multiple deep learning models, including Custom CNN, MobileNetV2, and ResNet50.
- Selected MobileNetV2 as the final deployment model based on its performance and computational efficiency.
- Developed a responsive web application using React and Flask.

- Integrated the trained deep learning model with the web application for real-time prediction.
- Implemented user authentication and profile management features.
- Successfully tested and validated all application modules.
- Deployed a functional AI-assisted Pneumonia Detection System.

These achievements demonstrate the successful completion of both the technical and functional objectives of the internship project.

10.3 Future Scope

Although the developed Pneumonia Detection System performs effectively, several enhancements can be incorporated in future versions to improve its capabilities and expand its practical applications.

Multi-Class Disease Detection

The current system performs binary classification between Normal and Pneumonia. Future versions can be extended to detect multiple chest diseases such as:

- Tuberculosis
- COVID-19
- Lung Cancer
- Pulmonary Fibrosis
- Pleural Effusion

This would make the application more versatile for clinical use.

Larger and More Diverse Dataset

Training the model using larger datasets collected from multiple hospitals and healthcare institutions would improve model generalization and reduce prediction bias.

A more diverse dataset would enable the model to perform better on real-world clinical images.

Explainable Artificial Intelligence (XAI)

Future implementations can incorporate Explainable AI techniques such as:

- Grad-CAM
- Heatmaps
- Attention Maps

These techniques would allow users and healthcare professionals to visualize the regions of the chest X-ray that influenced the model's prediction, improving transparency and trust.

Cloud Deployment

The application can be deployed on cloud platforms such as:

- AWS
- Microsoft Azure
- Google Cloud Platform

Cloud deployment would improve scalability, availability, and accessibility while supporting a larger number of concurrent users.

Mobile Application

A dedicated Android or iOS application can be developed to provide mobile access to the Pneumonia Detection System, enabling healthcare professionals to perform predictions directly from smartphones or tablets.

Hospital Integration

The application can be integrated with existing Hospital Information Systems (HIS) and Electronic Health Record (EHR) systems, enabling seamless data exchange and improved clinical workflows.

Performance Optimization

Future research may investigate newer deep learning architectures, such as EfficientNet or Vision Transformers (ViTs), to further improve prediction accuracy and inference speed.

Model optimization techniques such as pruning and quantization can also be explored to reduce computational requirements.

Continuous Learning

Future versions of the system can incorporate mechanisms for continuous model updates using newly collected anonymized chest X-ray images, enabling the model to adapt to evolving clinical data and improve prediction performance over time.

10.4 Objectives Achieved

The internship concluded with a 100% success rate against its primary objectives:

The Pneumonia detection model was successfully scaled to 1000 images, proving empirically that volume eliminates overfitting, culminating in a highly accurate MobileNetV2 architecture.

A full-stack, responsive web dashboard and decoupled Flask REST API were successfully architected, integrated, and deployed.

The computer vision module successfully utilized MediaPipe and OpenCV to achieve real-time 3D landmark extraction. Complex Euclidean heuristics and temporal debouncing buffers were successfully implemented, ensuring zero-latency, highly stable gesture classification (Open Palm vs. Closed Fist).

10.5 Learning Outcomes

The internship provided valuable practical exposure to Artificial Intelligence, Deep Learning, and Full-Stack Web Development. Through this project, the following skills and knowledge were developed:

Technical Skills

- Deep Learning using TensorFlow and Keras.
- Transfer Learning with MobileNetV2 and ResNet50.
- Medical image preprocessing using OpenCV and Pillow.

- Frontend development using React.
- Backend development using Flask.
- REST API development and integration.
- Model deployment and integration.
- Git and GitHub version control.

Professional Skills

- Problem-solving and analytical thinking.
- Software development lifecycle implementation.
- Project planning and execution.
- Technical documentation and report writing.
- Testing and debugging.
- Team collaboration and communication.

These learning outcomes have strengthened the understanding of both Artificial Intelligence and software engineering practices while providing hands-on experience in developing real-world healthcare applications.

10.4 Future Enhancements

While the current systems are highly functional, they provide a strong foundation for future technological expansion.

Pneumonia System Future Scope: The immediate next step is the integration of Explainable AI (XAI) techniques, such as Grad-CAM. This would allow the system to generate heatmaps overlaid on the X-ray, visually highlighting to the doctor the exact regions of the lung that caused the neural network to predict pneumonia, vastly increasing clinical trust.

Behavior Tracking Future Scope: The gesture module will be expanded to support multi-hand tracking. Furthermore, the underlying MediaPipe architecture can be expanded to incorporate facial mesh recognition and full-body holistic pose estimation, ultimately creating a comprehensive, enterprise-grade behavioral analytics suite.

References

11. [1] F. Chollet, *Deep Learning with Python*, 2nd ed. Shelter Island, NY, USA: Manning Publications, 2021.
12. [2] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L. C. Chen, "MobileNetV2: Inverted Residuals and Linear Bottlenecks," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.
13. [3] C. Lugaresi, J. Tang, H. Nash, C. McClanahan, E. Uboweja, M. Hays, F. Zhang, C. Chang, M. Yong, J. Lee, and W. Chang, "MediaPipe: A Framework for Building Perception Pipelines," *arXiv preprint arXiv:1906.08172*, 2019.
14. [4] M. Grinberg, *Flask Web Development: Developing Web Applications with Python*. Sebastopol, CA, USA: O'Reilly Media, 2018.
15. [5] Mozilla Developer Network (MDN), "Using Fetch API." [Online]. Available: https://developer.mozilla.org/en-US/docs/Web/API/Fetch_API/Using_Fetch. [Accessed: July 2026].
16. [6] D. S. Kermany, M. Goldbaum, W. Cai, C. C. Valentim, H. Liang, S. L. Baxter, *et al.*, "Identifying Medical Diagnoses and Treatable Diseases by Image-Based Deep Learning," *Cell*, vol. 172, no. 5, pp. 1122–1131, 2018.
17. [7] freeCodeCamp.org, "TensorFlow 2.0 Complete Course – Python Deep Learning Tutorial," YouTube. [Online]. Available: <https://www.youtube.com/@freecodecamp>
18. [8] TensorFlow, "TensorFlow Official YouTube Channel," YouTube. [Online]. Available: <https://www.youtube.com/@TensorFlow>
19. [9] CodeBasics, "Deep Learning Tutorials using TensorFlow & Keras," YouTube. [Online]. Available: <https://www.youtube.com/@codebasics>
20. [10] Nicholas Renotte, "TensorFlow and Computer Vision Tutorials," YouTube. [Online]. Available: <https://www.youtube.com/@NicholasRenotte>
21. [11] CampusX, "Deep Learning with TensorFlow and Keras," YouTube. [Online]. Available: <https://www.youtube.com/@campusx-official>
22. [12] Krish Naik, "Deep Learning and Machine Learning Tutorials," YouTube. [Online]. Available: <https://www.youtube.com/@krishnaik06>
23. [13] OpenCV, "OpenCV Official YouTube Channel," YouTube. [Online]. Available: <https://www.youtube.com/@OpenCV>
24. [14] Tech With Tim, "Flask Web Development Tutorials," YouTube. [Online]. Available: <https://www.youtube.com/@TechWithTim>
25. [15] Programming with Mosh, "React JS Full Course," YouTube. [Online]. Available: <https://www.youtube.com/@programmingwithmosh>
26. [16] Traversy Media, "React.js & Flask Development Tutorials," YouTube. [Online]. Available: <https://www.youtube.com/@TraversyMedia>